

`xdc·spec` is a special-purpose language for expressing the logical structure of software content in terms of three higher-level constructs – *modules*, *interfaces*, and *packages*. Enforcing a clean separation between software specification and software implementation, an `xdc·spec` source file effectively serves as a programmatic contract between content suppliers (producers) and their clients (consumers). Incorporating familiar C constructs for defining client-visible constants, types, and functions, a specification expressed in `xdc·spec` often resembles a “cleaned-up” rendition of a legacy header file.

While comparable to an IDL (Interface Definition Language), the true power of `xdc·spec` lies in its ability to forge a *single* specification that serves *two* congruent programming domains: [1] a target-domain, where the specified content is implemented in C/C++ and executes on a suitable hardware platform; and [2] a meta-domain, where this same content is potentially configured to match specific system requirements *before* target execution begins. In support of the latter, `xdc·spec` works hand-in-hand with `xdc·script` – a general-purpose language based on industry-standard JavaScript – used by producers and consumers alike to implement host-based *meta-programs* that ultimately beget target executables.

lexical-elements

The lexical structure of `xdc·spec` closely tracks that of JavaScript, which in turn is patterned after familiar conventions originating in C. The following table summarizes the different sorts of lexical elements potentially found in an `xdc·spec` source file:

Element	Examples
Whitespace	<i>space, tab, newline</i>
single-line comment	<i>// text</i>
block comment	<i>/* text */</i>
C-style identifier	<i>main, int32, _private</i>
JavaScript number	<i>123, 0xff, 6.02e23</i>
JavaScript string	<i>"hello\n", 'abc'</i>

syntactic-elements

The following table summarizes the different sorts of syntactic elements appearing throughout the `xdc·spec` reference grammar:

Element	Examples
syntactic production	<i>requires-statement, unit-category</i>
language keyword	<i>package, module</i>
language identifier	<i>dirname, UnitName</i>
literal value	<i>number, string</i>
literal symbol	<i>[, }, ;</i>
documentation comment	<i>///! @xdoc</i>

optional term	<i>grammar-term</i> ?
alternative terms	<i>grammar-term</i> <i>grammar-term</i>
zero-or-more occurrences	<i>grammar-term</i> *
one-or-more occurrences	<i>grammar-term</i> +
group-of terms	[<i>grammar-term</i> <i>grammar-term</i>]

package-specification

Within the realm of `xdc.spec`, a package is a programmatic element that logically contains modules and interfaces – collectively termed *units* – within its scope. Besides introducing its own public name, a package specification may identify other named packages upon which the current package in some way depends. The specification may also declare its (in)compatibility with earlier versions of the same package, as well as further constrain the level of compatibility required of any dependent packages. Specially-formatted comments embedded in package-specification source file are processed by the `xdoc` utility when generating HTML documentation for a set of packages.

package-specification

`requires-statement`*

`///! @xdoc`

`package` `qualified-package-name` `compatibility-key`?

{

`unit-declaration-list`*

} [;]?

`///! @xdoc`

requires-statement

`requires` `qualified-package-name` `compatibility-key`? ;

qualified-package-name

`dirname` [`.dirname`]*

compatibility-key

[`number` [`,number`]*]

unit-declaration-list

`unit-category` `UnitName` [`,UnitName`]* ;

unit-category

`module` | `interface`

-
- A `qualified-package-name` should in general be globally-unique. Packages are located in a directory with a matching name found along the system package path.
 - A `package-specification` resides in a distinguished source file named `package.xdc`, found in the corresponding package directory.

- Each **unit-declaration-list** introduces individual modules and interfaces within the scope of the current package, all of which must be uniquely named. Units with the same name may only appear in different packages.
- Each **requires-statement** designates another package upon which the current package in some way depends. The requires-relation among packages cannot contain cycles.
- A **compatibility-key** is generally interpreted as a numeric array of the form `[m,s,r,p]` where `m` denotes major functionality, `s` denotes source level, `r` denotes specification radius, and `p` denotes a particular patch. A new release of a package is *source-compatible* with any predecessor in which `m` is the same, and is *binary-compatible* with that predecessor if `s` is the same as well.
- By convention, package names are composed of `lowercase` identifiers whereas units (modules or interfaces) are named with `TitleCase` identifiers.

unit-specification

A unit specification defines all client-visible programmatic features of a module or interface. While virtually identical vis-à-vis their `xdc-spec` syntax, semantically this pair of programmatic elements are almost opposites: a module is *concrete* and *closed*, comprising a public specification accompanied by a conforming implementation; an interface is *abstract* and *open*, comprising only a public specification which others may import and ultimately implement. In general, the specification of a module or interface can inherit features from exactly one other (interface) specification, which itself can inherit additional features in the same manner – as if the latter’s cumulative set of features had been directly defined with the unit specification of its inheritor.

The features defined within the scope of a unit partition themselves into two main groups: *module-wide* features, which are associated with a solitary programmatic object encapsulating the implementation of some concrete module; and *per-instance* features, which are associated with a family of programmatic objects individually created and manipulated by the same underlying module. Aside from auxiliary definitions of supporting constants and types (deemed module-wide for convenience), features with a meaningful run-time presence in the underlying implementation – assignable configuration parameters along with callable functions – can participate in either category. As with packages, special documentation comments can be associated with the module/interface as a whole as well as with any named feature defined within its scope.

In the general case, a single unit specification for a module or interface defines a presence in two congruent programming domains: [1] the *target*-domain, where specified features become accessible within executable programs (written in C/C++) running on some particular platform; and [2] the *meta*-domain, where specified features become accessible within hosted scripts (written in `xdc-script`) used to configure these very same target programs ahead of their execution. In some cases, the specification of a particular module or interface can be restricted to the meta-domain – useful when constructing configuration “facades” atop legacy content as well as when distributing host-based meta-content for use in a broader range of scripting contexts.

unit-specification

```
///! @xdoc
metaonly? unit-category UnitName
  inherited-interface?
  {
    module-wide-feature*
  [ instance:
    per-instance-feature*
  ]?
  } [;]?
///! @xdoc
```

inherited-interface

```
inherits qualified-unit-name
```

qualified-unit-name

```
[qualified-package-name .]? UnitName
```

module-wide-feature

```
///! @xdoc
[
  | auxiliary-definition
  | config-parameter
  | function-declaration
]
```

per-instance-feature

```
///! @xdoc
[
  | config-parameter
  | function-declaration
]
```

-
- A **qualified-unit-name** effectively extends the (already) globally-unique name of the unit's containing package. If the **qualified-package-name** prefix is absent, the current package name is presumed.
 - A **unit-specification** resides in a source file named *UnitName.xdc*, found in the directory of its containing package.
 - Each **module-wide-feature** or **per-instance-feature** introduces an individual feature with the scope of the current unit, all of which must be uniquely named. Features with the same name may only appear in different units.
 - An **inherited-interface** designates a *single* interface whose features are introduced within the scope of the current unit. Any module or interface can optionally inherit features from another interface, so long as the inherits-relation among all units remains acyclic.
 - Each unit manifests itself in both the target- and meta-programming domains, unless designated **metaonly**. As a rule, inheritors of metaonly interfaces must themselves be metaonly modules or interfaces.

auxiliary-definition

An `xdc-spec` auxiliary definition defines a (module-wide) constant or type, often supporting other module-wide or per-instance features defined in the same unit. These definitions are also used by clients who consume this unit, both within any modules or interfaces they may specify as well as within any target (or meta) content they may implement. For the most part, the syntax and semantics of each form of auxiliary definition is patterned after a familiar programmatic construct already found in C. Further semantic restrictions guarantee auxiliary definitions have a meaningful manifestation in the meta-domain as well as in the target-domain.

auxiliary-definition

`const` | `enum` | `extern` | `struct`

`const`

`const` *typed-declaration* = *initializer* ;

`enum`

```
enum EnumName {  
    enum-value [, /*! @xdoc  
    enum-value ]* /*! @xdoc  
}
```

`enum-value`

ENUMVAL_NAME [= *initializer*]?

`extern`

`extern` *typed-declaration* = *symbolName* ;

`struct`

```
struct StructName {  
    struct-field*  
}
```

`struct-field`

typed-declaration ; */*! @xdoc*

`typedef`

`typedef` *typed-declaration* ;

-
- A `const` is restricted to numeric types, either standard or else enumerated. Its persistent value is defined by a statically-evaluated `initializer` consistent with its `typed-declaration`.
 - An `enum` is a new numeric type that ranges over a finite set of named values. As in C, integer values beginning with 0 are assigned to each successive `enum-value` (unless altered by an explicit numeric `initializer`). Each named value is actually resident in the scope of the containing module or interface, and hence must be unique among all features defined in this unit.

- An **extern** is a special form of constant that effectively aliases an external program symbol naming a C-language function or variable. Its **typed-declaration** is restricted to standard C types, including arbitrary pointer-types.
- A **struct** defines a new aggregate type comprising a set of assignable fields of any type. Each **struct-field** must be uniquely named within the scope of the enclosing struct.
- A **typedef** effectively defines a *synonym* for the type specified in its **typed-declaration**, rather than a new type per se. As in C, typedef names can appear in other typed declarations not unlike previously-defined enum or struct names.
- By convention: `UPPER_CASE` identifiers are used to name a const, enum value, or extern; `TitleCase` identifiers are used to name an enum/struct types as well as in typedefs; and `camelCase` identifiers are used to name struct fields.

config-parameter

A configuration parameter is a feature that behaves like a “property” of the underlying module or instance object – a readable (and sometimes writeable) variable of virtually any type. In the most general case, module-wide configuration parameters are assigned within the meta-domain and then become persistent constants within the target-domain; per-instance configuration parameters are likewise assignable within the meta-domain, but are limited to supporting run-time instance creation within the target-domain. Where appropriate, configuration parameters can be restricted to the meta-domain as well as designated readonly after initialization.

A configuration parameter inherited from some previously specified interface can itself be overridden – typically to (re-)define its initial value. A configuration parameter can also be finalized, effectively freezing its definition and precluding further overrides.

config-parameter

config-modifiers

readonly? **config** **typed-declaration** [**= initializer**]? ;

config-modifiers

final? **override?** **metaonly?**

-
- The optional **initializer** must yield a value consistent with the **typed-declaration**, using the rules of assignment-compatibility defined by `xdc-script` for the meta-domain. If no initializer is supplied, the configuration parameter starts out **undefined** in the meta-domain.
 - If **override** is specified among the **config-modifiers** for a configuration parameter, its **typed-declaration** as well as its use of **readonly** and **metaonly** must *exactly* match that of the inherited configuration parameter being overridden. Configuration parameters marked **final** cannot be overridden.
 - A **readonly** configuration parameter without an **initializer** can still be assigned a persistent value in the meta-domain, during construction of an underlying module/instance object.

- By convention, configuration parameters names are `camelCase` identifiers.

function-declaration

An `xdc-spec` function declaration generally stipulates the signature – argument and return types – of a callable routine implemented through a concrete module in either the target-domain or else (if so indicated) in the meta-domain; a target function also manifests itself in the meta-domain as an extern symbol of a function-pointer type derived from the stipulated signature. Following C++, default values can be specified for the last k arguments of an n -ary function, enabling the same routine to be called with as few as $n-k$ inputs; an untyped sequence of optional trailing arguments can be also specified using the familiar `...` notation. For those meta-domain functions wishing to adopt a more “weakly-typed” style supported (but not necessarily encouraged!) by the `xdc-script` language, their corresponding declaration in `xdc-spec` can just contain the names for each argument.

Like configuration parameters, a function declaration inherited from a previously-specified interface can be overridden – typically, to alter or extend the set of default argument values or else to allow optional trailing arguments. Note, though, that since interfaces are entirely abstract (void of any “default” implementation), inheritance of functions is limited to their client-visible specification; ultimately, it is concrete modules (or their delegates) that bear responsibility for implementing all functions directly or indirectly declared in their specification. With `xdc-spec` support for *Design-By-Contract* forthcoming, overriding inherited function declarations becomes an essential technique for weakening pre-conditions and strengthening post-conditions specified previously through executable expressions.

function-declaration

`typed-function-declaration` | `untyped-function-declaration`

typed-function-declaration

`function-modifiers`

`typed-declaration` (`typed-arguments?` [`,` `...`] `?`) ;

function-modifier

`final?` `override?` `metaonly?`

typed-arguments

`arg-declaration` [`,` `arg-declaration`] *

arg-declaration

`typed-declaration` [`= initializer`] ?

untyped-function-declaration

`function` `fxnName` (`untyped-arguments?`) ;

untyped-arguments

`argName` [`,` `argName`] *

- If **override** is specified among the **function-modifiers** for a function, its name and type signature as well as its use of **metaonly** must *exactly* match that of the inherited declaration being overridden. Functions marked **final** cannot be overridden.
- The optional **initializer** within an **arg-declaration** must yield a value consistent with the **typed-declaration**, using the rules of assignment-compatibility defined by **xdc-script** for the meta-domain.
- All argument names – whether **typed-arguments** or **untyped-arguments** – must be uniquely named on a per-function basis.
- An **untyped-function-declaration** is implicitly modified **final** and **metaonly**.
- By convention, function and argument names are `camelCase` identifiers.

typed-declaration

Generic declarations of typed identifiers patterned after the familiar (and sometimes awkward) syntax of C lie at the heart of virtually all **xdc-spec** feature definitions. These declarations stipulate a type name, either built-in or previously-defined, followed by what is conventionally termed a *declarator* – the name of the feature per se, optionally adorned with other syntactic elements. As in C, use of the *****, **[]**, and **()** operators within the declarator denotes new types such *pointer-to*(*t*), *array-of*(*t*), and *function-returning*(*t*) for some base type *t*; extra parentheses are typically used to bind the lower-precedent ***** operator to the declared name, especially when defining types of form *pointer-to*(*function-returning*(*t*)).

Beyond these familiar C constructs lifted from the target-domain – each given a corresponding meaning in the meta-domain – **xdc-spec** introduces additional base types as well as more specialized forms of the array: [1] the built-in type **any**, which subsumes all other types in the meta-domain; [2] the keywords **Module** or **Instance**, signifying an opaque type referencing a concrete module or instance object whose visible features are limited to those specified in the corresponding named unit; [3] the keyword **length** in conjunction with the **[]** operator, signifying the type *vector-of*(*t*) whose length can be altered in the meta-domain and retrieved in the target-domain; and [4] the keyword **string** in conjunction with the **[]** operator, signifying the type *map-into*(*t*) that effectively overlays direct access via string-valued keys on an underlying *vector-of*(*t*).

typed-declaration

[builtin-type | defined-type] declarator

builtin-type

standard-C-type | *extended-XDC-type*

defined-type

[qualified-unit-name .]? defined-type-name

defined-type-name

EnumName | *StructName* | *TypedefName* | **Instance** | **Module**

declarator

```
declared-name ?  
| * declarator  
| declarator [ [number | length | string]? ]  
| declarator ( argument-types ? [...] ? )  
| ( declarator )
```

declared-name

```
CONST_NAME  
| EXTERN_NAME  
| TypedefName  
| argName  
| configName  
| fieldName  
| fxnName
```

argument-types

```
typed-declaration [, typed-declaration ] *
```

-
- A **builtin-type** extends the familiar set of standard C types (`int`, `unsigned char`, etc.) with a set of capitalized equivalents (`Int`, `UChar`, etc.) as well as a handful of scalar types (`String`, `Ptr`, `Int32`, etc.) that promote further portability of C code. [Additional information on these extended types – both their definition in terms of standard C as well as their interpretation in the meta-domain – will be found elsewhere.]
 - A **defined-type** identifies a previously defined type, either in the module or interface designated by a valid **qualified-unit-name** or else in the current unit.
 - The **declared-name** at the heart of a **declarator** is only optional within **arguments-types**, typically used when declaring a type *pointer-to (function-returning(τ))*. C reference grammars usually refer to this syntactic construct as an *abstract-declarator*.

initializer

An initializer is an expression that denotes either a scalar or aggregate value, and whose elementary terms are manifest constants of known types. While certainly C-like in form and substance, xdc-spec scalar initializers are in fact statically-evaluated using JavaScript semantics (upon which xdc-script is based). Building on this foundation, xdc-spec aggregate initializers adopt standard JavaScript notation for denoting object and array values.

initializer

```
scalar-initializer | array-initializer | struct-initializer
```

scalar-initializer

literal
| defined-constant
| unary-op scalar-initializer
| scalar-initializer binary-op scalar-initializer
| scalar-initializer ? scalar-initializer : scalar-initializer
| (scalar initializer)

literal

number | *string* | **true** | **false** | **null** | **undefined**

defined-constant

[*qualified-unit-name* **.**]? *defined-constant-name*

defined-constant-name

CONST_NAME | *ENUMVAL_NAME*

unary-op

***** | **-** | **~** | **!**

binary-op

+ | **-** | ***** | **/** | **%** | **<<** | **>>** | **==** | **!=** | **<** | **<=** | **>** | **>=** | **&** | **|** | **^**

array-initializer

[] | [*initializer* [, *initializer*]* [,]?]

struct-initializer

{ } | { *field-initializer* [, *field-initializer*]* [,]? }

field-initializer

fieldName : *initializer*

-
- A **literal** *number* or *string* must conform to standard JavaScript, which also tracks standard C in this regard.
 - A **defined-constant** identifies a previously defined constant, either in the module or interface designated by a valid **qualified-unit-name** or else in the current unit.
 - The meaning and precedence of each **unary-op** and **binary-op** conforms to standard JavaScript, which likewise mirrors C. Note that JavaScript often overloads operators like + and < to accept strings as well as numbers.

// ! @xdoc

Special comments embedded within a specification source file are available for processing by independent tools used (say) for generation of HTML-based manuals or for interactive package browsing. These documentation comments are identified by an extra leading “bang” character – **// !** for single-line comments, **/* !** for the block variety – and can be juxtaposed with most named elements in the specification. Whenever multiple comments of either variety are associated with an individual specification element, their bodies are effectively concatenated into a single *documentation block* comprising one or more lines of text.

Markup of the form `@tag`, when present at the beginning of a line, further punctuates a documentation block into distinct sections comprising various styles of paragraphs. To avoid clutter in the source file due to excessive markup, most commentary can be written as “plain text” that follows some simple conventions to indicate (say) change-of-font or end-of-paragraph. Each documentation block generally comprises: [1] an untagged *summary* section, which is typically a “one-liner”; [2] an optional untagged *details* section, which may contain multiple paragraphs with additional information; and [3] a series of *tagged* sections, which further compartmentalize information about the associated specification element.

[Additional information on `@xdoc` – covering the current set of supported tags as well as textual hints – will be found elsewhere.]

```
//! @xdoc
    summary-section
    details-section?
    tagged-section*
summary-section
    comment-paragraph
details-section
    comment-paragraph+
tagged-section
    @sectTag(ident) comment-paragraph+
comment-paragraph
    [@p(style)]? textLines*
```